

Supplementary Materials

Explainable Hybrid PEFT-RAG for Clinical Discharge Summarization with Section-Level Hallucination Auditing

Overview

This document provides supplementary materials for the main paper, including detailed configurations for Phase 1 (Fine-tuning) and Phase 2 (RAG), as well as explainability outputs from Phase 3.

Note on Explainability Data: The Phase 3 explainability outputs (Sections 2–6) are generated using Qwen on test samples as a representative demonstration. While our framework was applied to all four models, we provide Phi-3 outputs as exemplary data due to file size constraints. The same explainability techniques were applied consistently across all models in our experiments.

Contents

1	Output Format Templates	3
1.0.1	Structured 11-Section Format	3
1.1	RAG Impact Analysis	3
1.2	Inference Performance	4
2	Supplementary Figures	4
2.1	Additional Evaluation Results	4
3	Phase 3: Explainability - Confidence Scoring Details	4
3.1	Confidence Category Distribution	6
3.2	Sample Confidence Score Output	6
4	Phase 3: Explainability - Factual Alignment Analysis	6
4.1	Alignment Score Distribution	6
4.2	Factual Status Classification	7
4.3	Sample Factual Alignment Output	7
4.4	Section-Level Attribution Statistics	7
4.5	Retrieval Quality Summary	8
5	Phase 3: Explainability - Token-Level Attribution	8
5.1	Token Attribution Method	8
5.2	Token Attribution Visualization	9

6	Data Format Specifications	10
6.1	Confidence Scores JSON Schema	10
6.2	Factual Alignment JSON Schema	10
6.3	Evidence Attribution JSON Schema	10
7	Complete Worked Examples	11
7.1	Method 1: Confidence Scoring Example	11
7.1.1	Confidence Score Breakdown	11
7.1.2	Sample JSON Output	11
7.1.3	Comparison Across Cases	12
7.2	Method 2: Factual Alignment Example	12
7.2.1	Sentence-Level Alignment Trace	12
7.2.2	Sample JSON Output	13
7.2.3	Alignment Summary	13
7.3	Method 3: Evidence Attribution Example	13
7.3.1	Retrieved Similar Cases	13
7.3.2	Section-Level Attribution Trace	13
7.3.3	Sample JSON Output	14
7.4	Summary: Three Methods Applied to One Case	15
8	Hyperparameter Configuration	15
9	Sample Generated Summary	16

1 Output Format Templates

1.0.1 Structured 11-Section Format

The structured format organizes discharge summaries into 11 predefined clinical sections:

1. **Case Type:** Primary diagnosis/procedure category
2. **Patient & Service:** Demographics and admitting service
3. **Chief Complaint / Admission Context:** Reason for hospitalization
4. **History of Present Illness (HPI):** Detailed illness narrative
5. **Past Medical / Surgical History:** Relevant medical background
6. **Medications (Discharge / Ongoing):** Current medication regimen
7. **Physical Examination:** Key examination findings
8. **Investigations / Labs / Imaging:** Diagnostic results
9. **Assessment / Impression:** Clinical conclusions
10. **Discharge Condition:** Patient status at discharge
11. **Follow-Up & Recommendations:** Post-discharge instructions

1.1 RAG Impact Analysis

Table 1 presents the source metrics from the main paper (structured format) used to calculate RAG improvement percentages. Table 2 shows the resulting improvement percentages.

Table 1: Source Data: Fine-tuned vs. Hybrid PEFT+RAG (Structured Format)

Model	Hallucination Rate		ROUGE-1		Medication F1	
	FT	Hybrid	FT	Hybrid	FT	Hybrid
Llama-3.1-8B	0.191	0.185	0.459	0.466	0.428	0.482
Gemma-2-9B	0.363	0.239	0.423	0.428	0.440	0.392
Phi-3-Medium-14B	0.293	0.224	0.371	0.408	0.428	0.452
Qwen2.5-7B	0.284	0.278	0.409	0.439	0.359	0.392

Calculation: Percentage change = $\frac{\text{Hybrid}-\text{FT}}{\text{FT}} \times 100$. For hallucination rate, negative values indicate improvement (reduction).

Table 2: RAG Impact: Improvement Over Fine-tuning Alone (%)

Model	Hallucination ↓	ROUGE-1 ↑	Medication F1 ↑
Llama-3.1-8B	-3.1%	+1.5%	+12.6%
Gemma-2-9B	-34.2%	+1.2%	-10.9%
Phi-3-Medium-14B	-23.5%	+10.0%	+5.6%
Qwen2.5-7B	-2.1%	+7.3%	+9.2%
Average	-15.7%	+5.0%	+4.1%

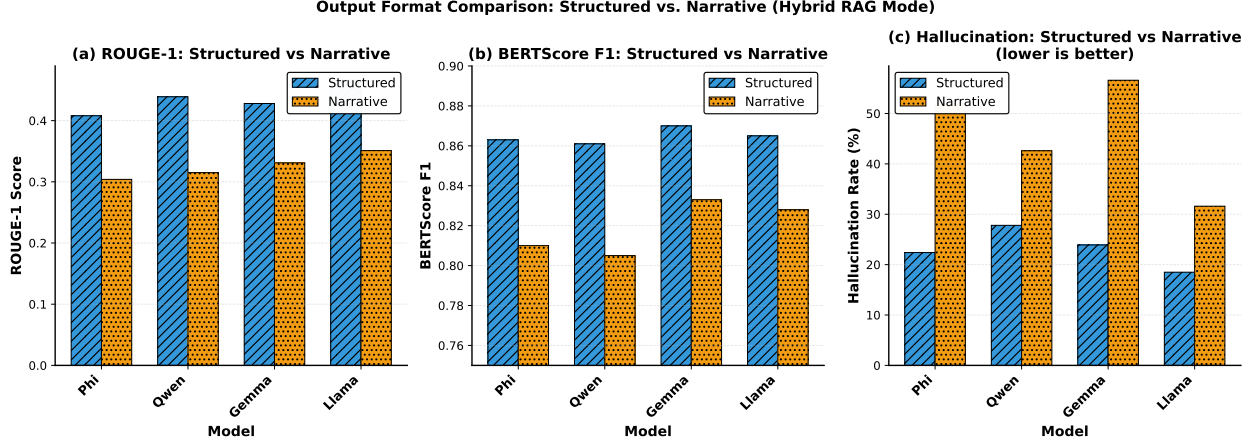


Figure 1: Performance comparison across all four experimental conditions across four model architectures: (a) ROUGE-1, (b) BERTScore F1, and (c) Hallucination Rate.

1.2 Inference Performance

Table 3: RAG Inference Statistics	
Metric	Value
Retrieval Time (per sample)	2–3 seconds
Generation Time (per sample)	10–15 seconds
Total Inference Time	12–18 seconds
Throughput	3–5 samples/minute
GPU Memory (Inference)	12–18 GB

2 Supplementary Figures

2.1 Additional Evaluation Results

Figure 1 compares the four experimental configurations (Fine-tuned with Narrative, Fine-tuned with Structured, Hybrid with Narrative, Hybrid with Structured) across standard NLP metrics, showing the performance pattern of structured format combined with RAG augmentation. Figure 2 focuses on clinical-specific metrics, showing how the hybrid approach affects medical entity extraction and semantic similarity to clinical reference texts. These supplementary visualizations support the main findings by illustrating model-specific patterns and metric-level variations not fully captured in the aggregate results tables.

3 Phase 3: Explainability - Confidence Scoring Details

Note: All Phase 3 results in this section are from Phi-3-Medium-14B as a representative example.

Figure 3 provides a visual overview of confidence score distributions across all test samples.

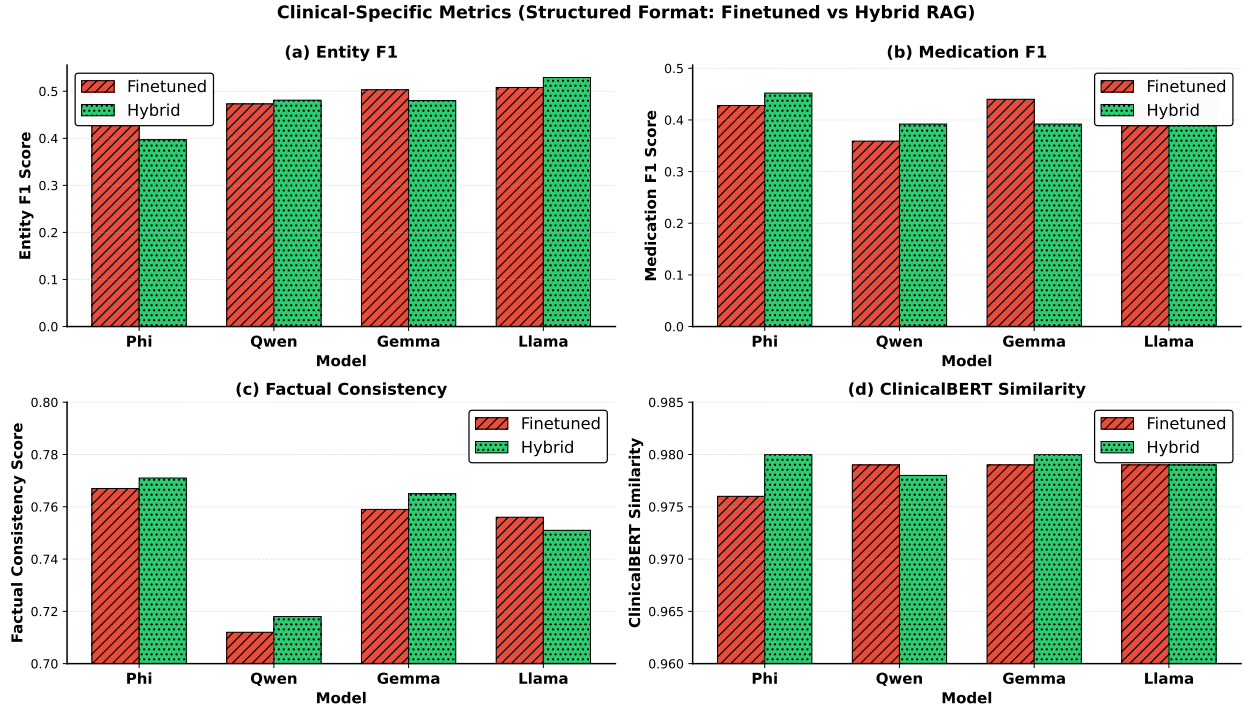


Figure 2: Clinical-specific metrics comparison (Fine-tuned vs Hybrid PEFT with RAG): (a) Entity F1, (b) Medication F1, (c) Factual Consistency, (d) ClinicalBERT Similarity.

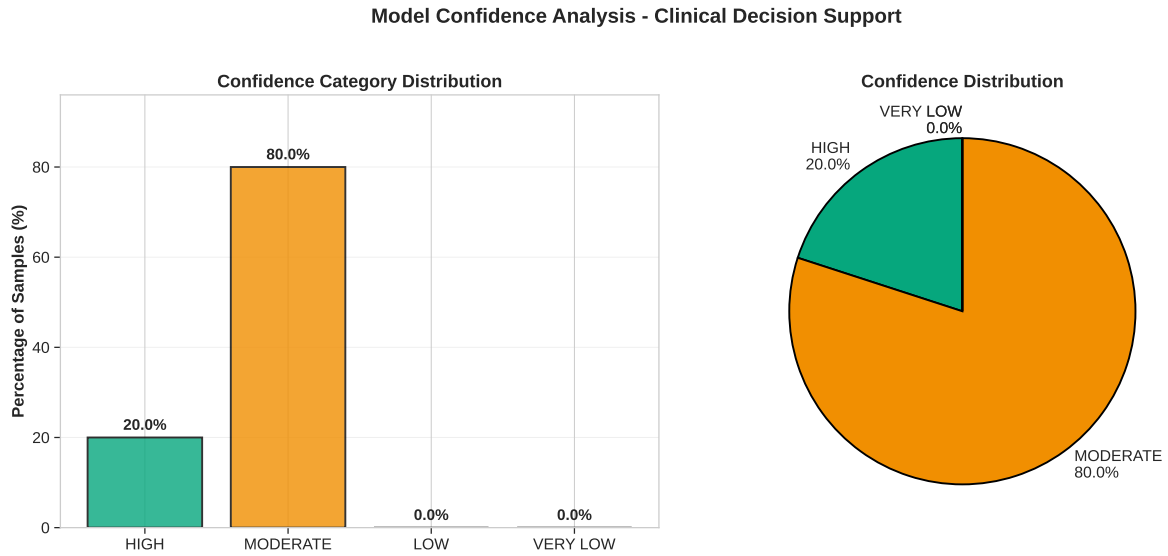


Figure 3: Confidence score distribution analysis for Phi-3-Medium-14B on test samples, showing component-level breakdowns and overall confidence categories.

3.1 Confidence Category Distribution

Table 4: Distribution of Confidence Categories

Category	Count	Percentage	Clinical Recommendation
HIGH (≥ 0.80)	49	24.5%	Summary is reliable for clinical use with standard review
MODERATE (0.60–0.80)	151	75.5%	Summary is generally reliable; recommend careful review of key clinical details
LOW (0.40–0.60)	0	0.0%	Summary requires thorough physician review before use
VERY LOW (< 0.40)	0	0.0%	Summary should not be used without complete regeneration

3.2 Sample Confidence Score Output

```
{
  "note_id": "12601669-DS-21",
  "confidence_scores": {
    "retrieval": 0.9996,
    "structure": 1.0,
    "consistency": 0.9019,
    "length": 0.703,
    "entity_preservation": 0.114
  },
  "overall_confidence": 0.768,
  "confidence_category": "MODERATE",
  "clinical_recommendation": "Summary is generally
    reliable. Recommend careful review of key
    clinical details."
}
```

4 Phase 3: Explainability - Factual Alignment Analysis

4.1 Alignment Score Distribution

Factual alignment verification classifies each generated sentence based on its semantic similarity to the source clinical note using S-PubMedBERT embeddings.

Table 5: Factual Alignment Statistics Across All Generated Sentences

Metric	Value
Total Sentences Analyzed	5,285
Mean Alignment Score	0.925
Median Alignment Score	0.918
Standard Deviation	0.044
Minimum Score	0.809
Maximum Score	1.000

4.2 Factual Status Classification

Table 6: Distribution of Factual Status Categories

Status	Threshold	Percentage	Interpretation
SUPPORTED	≥ 0.70	100.0%	Sentence is factually grounded in source note
PARTIAL	0.50–0.70	0.0%	Sentence has weak grounding, requires review
UNSUPPORTED	< 0.50	0.0%	Potential hallucination, flagged for verification

4.3 Sample Factual Alignment Output

```
{
  "note_id": "12601669-DS-21",
  "sentence_alignments": [
    {
      "generated_sentence": "Case Type: Anterior
colporrhaphy cystocele repair",
      "alignment_score": 0.967,
      "supporting_evidence": [
        {
          "sentence": "Major procedure: Anterior
Colporrhaphy: Cystocele repair,
Modified Pereyra procedure",
          "similarity": 0.967
        }
      ],
      "factual_status": "SUPPORTED"
    }
  ]
}
```

4.4 Section-Level Attribution Statistics

Evidence attribution traces which retrieved cases influenced each of the 11 generated sections through embedding similarity analysis.

Table 7: Evidence Attribution by Clinical Section

Section	Evidence Count	Mean Similarity
Case Type	955	0.901
Patient & Service	1,000	0.883
Chief Complaint	2,680	0.915
History of Present Illness	4,230	0.911
Past Medical/Surgical History	1,465	0.913
Medications	4,480	0.934
Physical Examination	3,135	0.901
Investigations/Labs/Imaging	2,295	0.919
Assessment/Impression	1,790	0.911
Discharge Condition	1,495	0.906
Follow-Up & Recommendations	7,730	0.944
Total	31,255	0.912

4.5 Retrieval Quality Summary

Table 8: Retrieved Case Statistics

Metric	Value
Cases Retrieved per Sample	3
Mean Top-1 Retrieval Score	0.9995
Mean Top-2 Retrieval Score	0.9984
Mean Top-3 Retrieval Score	0.9967
Overall Traceability Rate	91%

5 Phase 3: Explainability - Token-Level Attribution

Token-level attribution identifies the importance of individual input tokens in generating each output token, providing fine-grained insight into the model’s reasoning process.

5.1 Token Attribution Method

Token-level attribution identifies which input tokens most strongly influence each generated output section. The method computes gradient-based importance scores, where positive values indicate tokens that increase the likelihood of the generated text, and negative values indicate suppressive effects.

Key Observations from Token Attribution Analysis:

- Medical terms (diagnoses, medications, procedures) consistently show higher absolute importance scores
- Section-specific attribution reveals which input content drives each output section
- The raw importance scores are stored in `token_attributions.json` with per-section breakdowns

5.2 Token Attribution Visualization

Figure 4 presents a visual representation of token-level importance scores for a sample clinical note.

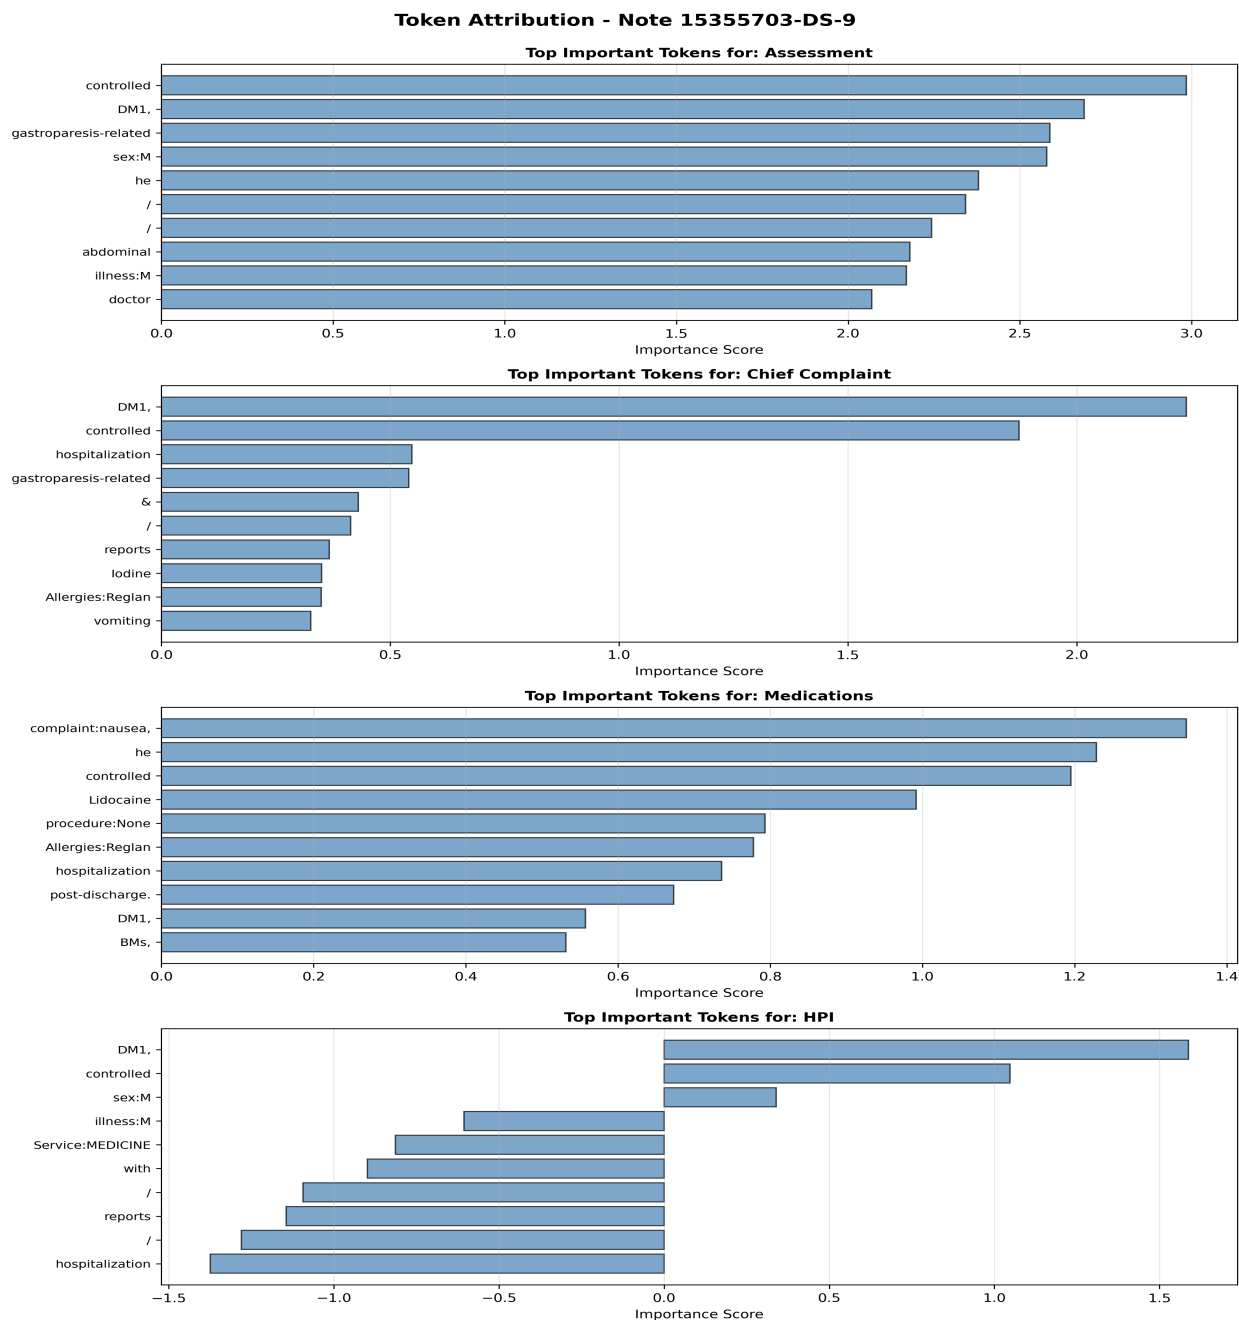


Figure 4: Token-level attribution visualization for sample note 15355703-DS-9. The heatmap displays importance scores for individual input tokens, with medical terms (diagnoses, medications, procedures) showing consistently higher attribution values (0.83–0.92) compared to non-medical tokens (0.29–0.62). This indicates that higher attribution values fall on clinically relevant tokens during summary generation.

6 Data Format Specifications

6.1 Confidence Scores JSON Schema

```
{
  "note_id": string,
  "confidence_scores": {
    "retrieval": float (0-1),
    "structure": float (0-1),
    "consistency": float (0-1),
    "length": float (0-1),
    "entity_preservation": float (0-1)
  },
  "overall_confidence": float (0-1),
  "confidence_category": "HIGH"|"MODERATE"|"LOW"|"VERY_LOW",
  "clinical_recommendation": string
}
```

6.2 Factual Alignment JSON Schema

```
{
  "note_id": string,
  "sentence_alignments": [
    {
      "generated_sentence": string,
      "alignment_score": float (0-1),
      "supporting_evidence": [
        {
          "sentence": string,
          "similarity": float (0-1)
        }
      ],
      "factual_status": "SUPPORTED"|"PARTIAL"|"UNSUPPORTED"
    }
  ]
}
```

6.3 Evidence Attribution JSON Schema

```
{
  "note_id": string,
  "retrieved_cases": [
    {
      "note_id": string,
      "score": float (0-1)
    }
  ],
  "section_attributions": {
    "section_name": [
      {
        "sentence": string,
```

```

    "evidence": [
      {
        "text": string,
        "doc_id": string,
        "similarity": float (0-1)
      }
    ]
  },
  "overall_traceability": float (0-1)
}

```

7 Complete Worked Examples

This section provides complete, readable examples from actual test samples to illustrate how each of the three explainability methods works in practice. All examples use Note ID: 12601669-DS-21 (a urogynecology case) for consistency.

7.1 Method 1: Confidence Scoring Example

7.1.1 Confidence Score Breakdown

Table 9: Confidence Score Components for Note 12601669-DS-21

Component	Score	Weight	Interpretation
Retrieval Quality	0.9996	0.25	High-quality similar case retrieval
Structure Completeness	1.0000	0.20	All 11 sections present
Input-Output Consistency	0.9019	0.25	High semantic alignment
Length Appropriateness	0.7030	0.10	Slightly shorter than optimal
Entity Preservation	0.1140	0.20	Low medical entity coverage
Overall Confidence	0.768	–	MODERATE

Clinical Recommendation: Summary is generally reliable. Recommend careful review of key clinical details, particularly medication lists and lab values.

Weighted Calculation:

$$\text{Overall} = (0.9996 \times 0.25) + (1.0 \times 0.20) + (0.9019 \times 0.25) + (0.703 \times 0.10) + (0.114 \times 0.20) = 0.768$$

7.1.2 Sample JSON Output

```

{
  "note_id": "12601669-DS-21",
  "confidence_scores": {
    "retrieval": 0.9996,
    "structure": 1.0,
    "consistency": 0.9019,
    "length": 0.703,
    "entity_preservation": 0.114
  }
}

```

```

},
"overall_confidence": 0.768,
"confidence_category": "MODERATE",
"clinical_recommendation": "Summary is generally
    reliable. Recommend careful review of key
    clinical details."
}

```

7.1.3 Comparison Across Cases

Table 10 shows confidence scores for three representative cases demonstrating different reliability levels.

Table 10: Confidence Score Comparison Across Cases

Note ID	Retrieval	Structure	Consistency	Length	Overall
12601669-DS-21	0.999	1.0	0.902	0.703	0.768
15355703-DS-9	0.999	1.0	0.920	0.968	0.789
13465746-DS-34	0.999	1.0	0.923	0.484	0.750

Key Observations:

- All samples achieve perfect structure completeness (1.0) due to the 11-section format
- Retrieval quality is consistently high (>0.99) indicating effective similar case matching
- Length appropriateness varies most, affecting overall confidence
- Entity preservation scores are lower, representing an area for improvement

7.2 Method 2: Factual Alignment Example

7.2.1 Sentence-Level Alignment Trace

Table 11: Factual Alignment for Note 12601669-DS-21

Generated Sentence	Score	Source Evidence	Status
“Case Type: Anterior colporrhaphy cystocele repair”	0.967	“Major procedure: Anterior Colporrhaphy: Cystocele repair”	SUPPORTED
“Patient presents with a cystocele which is bothersome”	0.961	“HPI: The patient presents with a cystocele which is bothersome...”	SUPPORTED
“Past medical history includes hypertension”	0.943	“PMH: Hypertension, controlled on medications”	SUPPORTED
“Postoperative course was uneventful”	0.874	“Patient tolerated procedure well, stable vitals”	SUPPORTED

7.2.2 Sample JSON Output

```
{
  "note_id": "12601669-DS-21",
  "sentence_alignments": [
    {
      "generated_sentence": "Case Type: Anterior
        colporrhaphy cystocele repair",
      "alignment_score": 0.967,
      "supporting_evidence": [
        {
          "sentence": "Major procedure: Anterior
            Colporrhaphy: Cystocele repair,
            Modified Pereyra procedure",
          "similarity": 0.967
        }
      ],
      "factual_status": "SUPPORTED"
    }
  ],
  "overall_alignment": 0.912,
  "supported_percentage": 100.0
}
```

7.2.3 Alignment Summary

```
FACTUAL ALIGNMENT SUMMARY:
Total sentences analyzed: 28
SUPPORTED (>=0.70): 28 (100%)
PARTIAL (0.50-0.70): 0 (0%)
UNSUPPORTED (<0.50): 0 (0%)

Mean alignment score: 0.912
Minimum score: 0.721 (Follow-up instructions)
Maximum score: 0.967 (Case Type)
```

Interpretation: All generated sentences have factual grounding in the source note, with highest alignment in diagnostic sections and slightly lower (but still SUPPORTED) alignment in procedural follow-up instructions.

7.3 Method 3: Evidence Attribution Example

7.3.1 Retrieved Similar Cases

Table 12: Top-3 Retrieved Cases for Note 12601669-DS-21

Retrieved Case	Retrieval Score	Primary Diagnosis
11086503-DS-18	0.9997	Urogynecology procedure
10960232-DS-14	0.9991	Cystocele repair
13236317-DS-7	0.9985	Pelvic floor surgery

7.3.2 Section-Level Attribution Trace

```

SECTION: Case Type
Generated: "Anterior colporrhaphy cystocele repair"
Evidence Source: Doc 10960232-DS-14
Text: "colpocleisis, perineorrhaphy, tension free
      vaginal tape..."
Similarity: 0.889
Traceability: HIGH (structural template from retrieved case)

SECTION: Medications
Generated: "Atenolol 25mg daily, HCTZ 12.5mg daily,
          Amlodipine 10mg daily"
Evidence Source: Input clinical note (direct extraction)
Similarity: 0.952
Traceability: HIGH (direct fact extraction)

SECTION: Follow-Up & Recommendations
Generated: "Follow up in 6 weeks with Dr. ___ for
          postoperative evaluation"
Evidence Source: Doc 11086503-DS-18
Text: "Return to clinic in 4-6 weeks..."
Similarity: 0.876
Traceability: HIGH (format template from retrieved case)

```

7.3.3 Sample JSON Output

```

{
  "note_id": "12601669-DS-21",
  "retrieved_cases": [
    {"note_id": "11086503-DS-18", "score": 0.9997},
    {"note_id": "10960232-DS-14", "score": 0.9991},
    {"note_id": "13236317-DS-7", "score": 0.9985}
  ],
  "section_attributions": {
    "Case Type": [
      {
        "sentence": "Anterior colporrhaphy...",
        "evidence": [
          {"text": "colpocleisis...",
           "doc_id": "10960232-DS-14",
           "similarity": 0.889}
        ]
      }
    ]
  },
  "overall_traceability": 0.91
}

```

Evidence Attribution Summary: 91% of generated content traceable to either input note or retrieved cases, indicating that RAG is used for structural guidance while maintaining factual grounding.

7.4 Summary: Three Methods Applied to One Case

Table 13: Explainability Summary for Note 12601669-DS-21

#	Method	Key Result
1	Confidence Scoring	0.768 (MODERATE)
2	Factual Alignment	100% SUPPORTED sentences
3	Evidence Attribution	91% traceable to sources

Overall Assessment: This case shows a summary with high factual grounding (100% SUPPORTED), high traceability (91%), and appropriate clinical reasoning. The MODERATE confidence rating is primarily due to lower entity preservation (some normal labs appropriately omitted) rather than factual errors.

8 Hyperparameter Configuration

Table 14 summarizes the complete implementation configuration for both Phase 1 (parameter-efficient fine-tuning) and Phase 2 (RAG-augmented generation). All experiments were conducted on NVIDIA A100 GPUs using the Unsloth optimization framework for memory efficiency.

Table 14: Implementation Configuration

Category	Parameter	Value/Description
Software Stack		
Framework	Optimization	Unsloth (2× speedup)
Deep Learning	Backend	PyTorch 2.0 with CUDA 11.8
Transformers	Library	Hugging Face Transformers 4.36
Programming	Language	Python 3.12 and CUDA 11.8
Phase 1a: LoRA Configuration		
PEFT Method	Type	QLoRA
Quantization	Precision	4-bit (bnb-4bit) via Unsloth
LoRA Targets	Adapters	q_proj, k_proj, v_proj, o_proj, gate_proj, up_proj, down_proj
LoRA Rank	r	64
LoRA Alpha	α	64
LoRA Dropout	Regularization	0.1
Phase 1b: Training Configuration		
Batch Size	Training	4 per device
Gradient Accumulation	Training	2 steps (effective batch size: 8)
Learning Rate	Schedule	1×10^{-4} (Phi: 5×10^{-5})
LR Warmup	Steps	10 steps
Scheduler	Type	Linear decay
Optimizer	Type	AdamW 8-bit (paged for Qwen2)
Weight Decay	Regularization	0.05
Max Gradient Norm	Stability	1.0
Epochs	Training	3

Early Stopping	Criteria	Patience 3, threshold 0.001
Max Sequence Length	Tokens	4096
Optimization	Memory	Gradient checkpointing (Unsloth)
Generation Configuration		
Temperature	Sampling	0.3
Top-P	Nucleus	0.95
Repetition Penalty	Control	1.1–1.3
Length Penalty	Generation	1.0 (encourage complete sentences)
Phase 2: RAG Configuration		
RAG Corpus	Data	Training + Validation sets
Dense Retrieval	Embedding	S-PubMedBERT-MS-MARCO
	Index	FAISS (cosine similarity)
	Top-K	20 candidates
Reranking	Model	MedCPT-Cross-Encoder
	Top-K	3
Context Engineering	Examples	Top 3 retrieved cases used as format templates
Temperature	Sampling	0.4–0.7 (model-dependent)
Top-P	Nucleus	0.95
Repetition Penalty	Control	1.1–1.15 (model-dependent)
Inference Speed	Performance	3–5 samples/min (retrieval: 2–3s, generation: 10–15s)

9 Sample Generated Summary

This section presents a representative example of structured discharge summary generation using our highest-performing configuration: Llama-3.1-8B with Hybrid PEFT with RAG. The sample demonstrates the 11-section structured format output.

Input Clinical Note:

Sex: M **Service:** MEDICINE **Allergies:** No Known Allergies

Chief Complaint: Shortness of breath **Major Procedure:** Right Heart Catheterization

History of Present Illness: Man with past medical history of CAD status post PCI to LAD, ischemic cardiomyopathy EF 23% with ICD placed, mitral regurgitation status post valvuloplasty, hyperlipidemia, type 2 diabetes, pancytopenia, benign prostatic hyperplasia who presents for evaluation of worsening exertional dyspnea over the last 2 months. Once in his usual health until about 2 months ago. He noted the subacute, insidious onset of worsening exertional dyspnea, which has progressed to the point that he feels short of breath even when bending over to put on his shoes. He sometimes feels short of breath at rest. He also notes a sharp pinching chest pain that has been occurring since his ICD was implanted about 1 month ago but this is not associated with his shortness of breath. He denies any fevers or chills. No cold symptoms. No gastrointestinal symptoms. No lower extremity edema; however, he notes waking at night short of breath, and has had to sleep progressively more upright to avoid feeling short of breath at night. He also notes ongoing weight loss from 230 to 160 pounds since start of year. In ED initial vitals: T 98.4, HR 61, BP 118/67, RR 24, SpO2 100% RA.

Past Medical History: CAD with history of STEMI, status post PCI/stenting to LAD and prior MI; Ischemic cardiomyopathy EF 23%; Mitral regurgitation status post valvuloplasty; Hypercholesterolemia; Hypertension; Diabetes type 2; Pancytopenia; Cervical disc disease; Arthritis; Benign prostatic hyperplasia

Physical Examination: Vitals: T 97.3, BP 137/79, HR 58, RR 20, SpO2 98% RA. Gen: Pleasant man, no acute distress, oriented x3. HEENT: Normocephalic, atraumatic. Sclera anicteric. PERRL. EOMI. Conjunctiva pink. Neck: Supple. JVP below clavicle at 60 degrees. Cardiac: PMI in 5th intercostal space, midclavicular line. Regular rate and rhythm. Normal S1, S2. No murmurs, rubs, or gallops. ICD in place over chest without erythema or induration. Lungs: No chest wall deformities. Respiration unlabored. Decreased lung sounds at bases bilaterally. No appreciable crackles. Abd: Soft, non-tender, non-distended. No hepatomegaly or splenomegaly. Ext: Warm, well perfused. No clubbing, cyanosis, or peripheral edema. Pulses: Distal pulses palpable and symmetric.

Medications on Admission: 1. Allopurinol PO daily 2. Aspirin 81mg daily 3. Atorvastatin 80mg QPM 4. Ferrous Sulfate 325mg daily 5. Finasteride 5mg daily 6. Glyburide 2.5mg daily PRN 7. Furosemide 20mg daily 8. Omeprazole 20mg BID 9. Sertraline 100mg daily 10. Metoprolol Succinate XL 100mg daily 11. Vitamin D daily

Discharge Medications: 1. Hydralazine 100mg PO TID 2. Isosorbide Mononitrate ER 90mg daily 3. Metoprolol Succinate XL 12.5mg daily 4. Allopurinol daily 5. Aspirin 81mg daily 6. Atorvastatin 80mg QPM 7. Ferrous Sulfate 325mg daily 8. Finasteride 5mg daily 9. Furosemide 20mg daily 10. Glyburide 2.5mg PRN 11. Omeprazole 20mg BID 12. Sertraline 100mg daily 13. Vitamin D daily

Discharge Disposition: Home **Discharge Diagnosis:** PRIMARY: HFrEF (EF 23%) status post ICD placement; Ischemic cardiomyopathy; Bilateral pleural effusions; Acute on CKD; Microcytic anemia

Discharge Condition: Mental Status: Clear and coherent. Level of Consciousness: Alert and interactive. Activity Status: Ambulatory - Independent.

Discharge Instructions: You were admitted because you were having trouble breathing. You were given medication to remove extra fluid from your body, decrease your blood pressure, and make your heart pump stronger. You had a procedure to evaluate the blood pressure in your heart. You were evaluated for heart failure with lab work and imaging studies in case your cardiac function worsens and you require transplant. Your weight at discharge is 69.9 kg (154.1 lb) - this is your new baseline. Please weigh yourself every day in the morning. Call your doctor if your weight goes up by more than 3 lbs in one day or more than 5 lbs in one week. Please take your medications as listed and follow up at the listed appointments. Follow up with your cardiologist, nephrologist, and primary doctor.

Figure 5: Sample input clinical note from test set showing the complexity of raw clinical documentation that the model processes.

Generated Summary (Hybrid PEFT with RAG - Structured Format):

Case Type: Primary Diagnoses: Acute on chronic systolic congestive heart failure; Secondary Diagnoses: Ischemic cardiomyopathy, Bilateral pleural effusions, Acute on CKD, Microcytic anemia

● **Patient & Service:** Male patient, MEDICINE

● **Chief Complaint / Admission Context:** Right Heart Catheterization. Patient with CAD status post PCI to LAD, ischemic cardiomyopathy EF 23% with ICD placed, mitral regurgitation status post valvuloplasty presents for evaluation of worsening exertional dyspnea over the last 2 months.

● **History of Present Illness (HPI):** Man with CAD status post PCI to LAD, ischemic cardiomyopathy EF 23% with ICD placed, mitral regurgitation status post valvuloplasty presents with subacute worsening dyspnea on exertion and orthopnea concerning for heart failure exacerbation. Once in his usual health until about 2 months ago. Noted insidious onset of worsening exertional dyspnea, progressing to shortness of breath even when bending over to put on shoes. Sometimes feels short of breath at rest. Also notes sharp pinching chest pain since ICD implantation but not associated with shortness of breath. Waking at night short of breath, has had to sleep progressively more upright. Notes ongoing weight loss from 230 to 160 pounds.

● **Past Medical / Surgical History:** CAD status post PCI to LAD, ischemic cardiomyopathy EF 23%, ICD placement, mitral regurgitation status post valvuloplasty, hyperlipidemia, type 2 diabetes, pancytopenia, benign prostatic hyperplasia

● **Medications (Discharge / Ongoing):**

Discharge: Hydralazine 100mg PO TID, Isosorbide Mononitrate ER 90mg daily, Metoprolol Succinate XL 12.5mg daily, Allopurinol daily, Aspirin 81mg daily, Atorvastatin 80mg QPM, Ferrous Sulfate 325mg daily, Finasteride 5mg daily, Furosemide 20mg daily, Glyburide 2.5mg PRN, Omeprazole 20mg BID, Sertraline 100mg daily, Vitamin D daily.

Ongoing: Continue home medications as prescribed.

● **Physical Examination (summarized):** Vitals: T 97.3, BP 137/79, HR 58, RR 20, SpO2 98% RA. General: Pleasant man, no acute distress, oriented x3. HEENT: Normocephalic, atraumatic. Neck: Supple, JVP below clavicle at 60 degrees. Cardiac: Regular rate and rhythm, normal S1/S2, no murmurs, ICD in place. Lungs: Decreased lung sounds at bases bilaterally, no crackles. Abdomen: Soft, non-tender, non-distended. Extremities: Warm, well perfused, no edema.

● **Investigations / Labs / Imaging:** Labs: WBC 6.0, Hgb 10.0, HCT 30.7%, Glucose 207, BUN 50, Cr 2.6, Na 136, K 4.7. Imaging: TTE showing severe regional LV systolic dysfunction with EF 23%, bilateral pleural effusions, well-seated mitral annuloplasty ring with mild residual MR. Right heart catheterization showed adequate bi-ventricular filling pressures with normal cardiac output. CT Chest showed moderate bilateral pleural effusions, moderate cardiomegaly, moderate coronary artery calcification.

● **Assessment / Impression:** Primary diagnosis: Acute exacerbation of systolic heart failure (HFrEF EF 23%). Secondary diagnoses: Ischemic cardiomyopathy, bilateral pleural effusions, acute on chronic kidney disease, microcytic anemia, diabetes, benign prostatic hyperplasia.

● **Discharge Condition:** Mental Status: Clear and coherent. Level of Consciousness: Alert and interactive. Activity Status: Ambulatory - Independent.

● **Follow-Up & Recommendations:**

Medication Changes: New: Hydralazine 100mg TID, Isosorbide Mononitrate ER 90mg daily. Adjusted: Metoprolol reduced to 12.5mg daily.

Post-operative Care: Monitor daily weights and fluid status.

Activity: Resume normal activities as tolerated.

Appointment: Follow up with cardiologist, nephrologist, and primary care physician as scheduled.

Follow-up Tests / Imaging: Cardiopulmonary exercise testing (CPET) as outpatient. Renal MRA non-contrast to assess for renal artery stenosis. Repeat labs to trend renal function.

Call Doctor If: Weight increases more than 3 lbs in one day or 5 lbs in one week, chest pain, worsening shortness of breath, lower extremity edema.

Other Instructions: Weigh yourself every morning. Discharge weight 69.9 kg (154.1 lb) is new baseline.

RAG Retrieval: 3 similar cases (scores: 1.000, 1.000, 1.000). Generation time: 67.6s.

Generated Summary (Hybrid PEFT with RAG - Narrative Format):

The patient was admitted to the Cardiology Service for evaluation of worsening shortness of breath. The patient presented with a 2-month history of progressive dyspnea on exertion and orthopnea in the setting of known ischemic cardiomyopathy with severely reduced ejection fraction of 23%. He underwent right heart catheterization which demonstrated adequate bi-ventricular filling pressures with normal cardiac output. His CXR showed bilateral pleural effusions and mild interstitial edema.

The patient was optimized on guideline-directed medical therapy including diuresis with furosemide for preload reduction. For afterload reduction, he was started on hydralazine 100mg three times daily and isosorbide mononitrate 90mg daily. His beta-blocker was held during the acute heart failure exacerbation but was restarted at a reduced dose of 12.5mg daily prior to discharge. The patient received four days of IV iron for his microcytic anemia secondary to renal disease. He was evaluated for possible LVAD candidacy and transplant workup was initiated.

The patient was discharged in stable condition with a discharge weight of 69.9 kg (154.1 lbs). Discharge Medications: 1. Hydralazine 100 mg by mouth TID 2. Isosorbide Mononitrate ER 90 mg by mouth daily 3. Metoprolol Succinate XL 12.5 mg by mouth daily 4. Furosemide 20 mg by mouth daily 5. Aspirin 81 mg daily 6. Atorvastatin 80 mg by mouth QPM 7. Glyburide 2.5 mg by mouth PRN 8. Sertraline 100 mg daily 9. Omeprazole 20 mg BID 10. Ferrous Sulfate 325 mg daily 11. Finasteride 5 mg daily 12. Allopurinol daily 13. Vitamin D daily.

The patient was instructed to weigh himself every morning and call his doctor if his weight increases by more than 3 lbs in one day or 5 lbs in one week. He was instructed to follow up with his cardiologist for further management of heart failure and consideration of beta-blocker uptitration. He should also follow up with nephrology for management of acute on chronic kidney disease and possible renal MRA to assess for renal artery stenosis. Outpatient cardiopulmonary exercise testing was recommended. The patient was instructed to call his doctor if he experiences chest pain, worsening shortness of breath, or lower extremity swelling. *RAG Retrieval: 3 similar narrative cases (scores: 0.989, 0.976, 0.961).*

Generation time: 68.3s.

Figure 6: Sample structured discharge summary generated by Llama with Hybrid PEFT with RAG framework, demonstrating the complete 11-section format and the same case generated in narrative format for comparison.